

Haoyu Dong

✉ hd2573@columbia.edu

📞 1-646-207-6669

🔗 kozmojor.github.io/haoyudong.github.io/

Education

Columbia University

MS in Electrical Engineering

SEP 2024 – DEC 2025 (EXPECTED)

GPA: 4.07/4.0

- **Awards:** 2025 Spring MS Honors Students as top 10% of the program.

Xi'an Jiaotong University

BE in Automation

SEP 2020 – JUL 2024

GPA: 3.51/4.3

- **Awards:** Excellent Student Scholarship, XJTU, Academic Year 2022-2023

Research Interest

Focus: Efficient use of **data** and **training** to build trustworthy **multimodal** intelligent systems for complex real-world tasks.

Keywords: Multimodal LLMs; Self-supervised/RL Post-training; Trustworthy Reasoning

Publications

Zhang, Z., **Dong, H.**, Pei, K. & Mao, C. (2025). *Cross-Modal Cycle Consistency Rewards Improve Multimodal Reasoning* ↗. (Under review; preprint available online.)

Zang, C., **Dong, H.**, Turkcan, M. K., Ghaderi, J., Zussman, G., & Kostic, Z. (2025). *Adaptive Data Collection for Robust LLM Post-training*. (Manuscript in Preparation.)

Research Experience

Cross Modality Self-play for Data-Free Training

Research Assistant | Advisor: Prof. Chengzhi Mao

NEW YORK, NY

AUG 2025 – NOW

- Proposed C³R, a cycle-consistency reward framework that resolves the modality gap in multimodal LLMs by replacing unstable majority-vote self-improvement with cross-modal self-verification.
- Designed four-path answer–query–answer cycles (T→T, T→I, I→T, I→I) to generate dense, label-free RL signals that correct cross-modal conflicts and stabilize multimodal reasoning.
- Achieved up to +7.6 pp accuracy and +10 pp consistency gains over Qwen-VL baselines and voting-based methods across six multimodal benchmarks including ScienceQA, ChartQA, MathVista, and VWA.
- Submitted as a full paper to CVPR 2026 and currently under review.

Adaptive Data Sampling in LLM Post-training

Research Assistant | AIDL Lab, Columbia University | Advisor: Prof. Zoran Kostic

NEW YORK, NY

FEB 2025 – AUG 2025

- Introduced a source-aware adaptive sampler for LLM that optimizes LLM post-training efficiency in online scenario.
- Formulated sampling as an adversarial multi-armed bandit; estimate per-source utility from training dynamics; update via UCB/EXP3 bandit under a fixed sampling budget; evaluate on GSM8K math reasoning against uniform sampling baselines.
- Obtained +5 pp accuracy on math reasoning and reached baseline performance with 10% fewer sampling budgets.

Enhanced Visual Reasoning Framework for Verifiable Language Model VQA

Leader | Advisor: Prof. Zoran Kostic

NEW YORK, NY

JAN 2025 – APR 2025

- Designed a visual reasoning framework with agentic consensus to reduce hallucination & inconsistency in knowledge-based VQA.
- Evidence-dense captioning (multi-view signals distilled into a single, entity/attribute-grounded scene description) + K-agent self-consistency (majority vote with rationale-consistency checks); evaluated on A-OKVQA / OK-VQA against strong CoT baselines.
- Achieved +13 pp overall accuracy improvement over vanilla CoT baselines; ablation shows +9 pp from evidence-dense captioning and +4 pp from agentic voting.

Linear-Time Rotary-Positional Attention for Efficient Vision Transformers

Leader | Advisor: Prof. Krzysztof Choromanski

NEW YORK, NY

SEP 2024 – MAR 2025

- Introduced RoPerformer, a ViT with rotary-linear attention to remove the quadratic attention bottleneck in high-resolution image classification while preserving relative positional bias.
- RoPerformer-ViT on ImageNet-1K; ablate projection rank r and 1D vs 2D RoPE; baselines: vanilla ViT under identical training.
- Cut step-time by 30–40% and peak VRAM by 25–35% at 384–512 px, holds Top-1 parity at 224 (± 0.2 pp), and yields

+0.6–1.0 pp Top-1 at higher resolutions; enables 1.6–2.0 \times larger batch under the same memory.

Academic Service

Conference Reviewer — *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*

2026

- Reviewer for CVPR 2026.

Teaching & Academic Engagement

Research Assistant — *CM Lab, Rutgers University*

AUG 2025 – JAN 2026

- Led WEBGUARD ARENA for adversarial self-play Web agent safety, and contributed to C³R cross-modal self-play with cycle-consistency rewards for label-free multimodal post-training.

Teaching Assistant — *Electrical Engineering, Columbia University*

SEP 2025 – DEC 2025

- Contributed to a graduate-level Deep Learning course. Assisted with assignments design & final coding projects assistance.

Research Assistant — *AIDL Lab, Columbia University*

JAN 2025 – AUG 2025

- Contributed to the project on data collection for LLM post-training. Worked under NSF traffic project.

Skills

Languages: Python, MATLAB, Shell, C, L^AT_EX, Markdown, HTML

Libraries & Tools: PyTorch, TensorFlow, Hugging Face Transformers, Verl, Git, LLM APIs.